

Data, Estimation and Inference Report

Ben Ellis

This report considers data from a sensor measuring tide height. The sensor often fails to transmit readings due to severe weather, and hence we here try to interpolate from the data available.

The readings from the sensor and the true tide heights over the period considered are shown in Figure 4 in the Appendix.

This data exhibits a few interesting patterns:

- As we would expect, the tide height is periodic with a period of roughly 12 hours.
- There are a few notable sections where data are entirely missing.
- Although data are missing, observed readings show little, and uniform, noise.

The regular structure of the data suggests that we could use Gaussian Processes to interpolate.

Gaussian Processes

A Gaussian process $f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), K(\mathbf{x}, \mathbf{x}))$ is a collection of random variables, any finite collection of which is Gaussian distributed. m is a function which gives the value of the mean of the Gaussian at its input, and K is a kernel function which given inputs \mathbf{x}_i and \mathbf{x}_j computes the entry in the covariance matrix Σ_{ij} . Hence for any point \mathbf{x}_\star the Gaussian process induces a Gaussian distribution of the possible function values $f(\mathbf{x}_\star)$. This is shown for a simple linear case in Figure 3 in the Appendix. More than that though, the Gaussian process also gives a joint Gaussian distribution for any discrete set of points X_\star , meaning that this allows us to control the variation in the function between points by specifying the kernel function.

Inference in Gaussian Processes

To perform inference, we must use the Gaussian Process (GP) to predict the function values \mathbf{f}_\star at some set of points X_\star given noisy observations \mathbf{y} at a set of points X . We assume that a single observation y_i is affected by IID Gaussian noise and hence the observations \mathbf{y} and function values \mathbf{f}_\star are jointly Gaussian.

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_\star \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} K(X, X) + \sigma^2 I & K(X, X_\star) \\ K(X_\star, X) & K(X_\star, X_\star) \end{bmatrix} \right)$$

where we have assumed that the mean is $\mathbf{0}$. In this case we normalised the data to zero mean and unit variance to achieve this. The identity matrix is used here because of the IID assumption about our noise. Since the conditionals of a Gaussian are also Gaussian, this gives us the posterior distribution

$$\begin{aligned} p(\mathbf{f}_\star | \mathbf{y}, X, X_\star) &= \mathcal{N}(\mathbf{m}_\star, \mathbf{K}_\star) \\ \mathbf{m}_\star &= K(X_\star, X)(K(X, X) + \sigma^2 I)^{-1} \mathbf{y} \\ \mathbf{K}_\star &= K(X_\star, X_\star) - K(X_\star, X)(K(X, X) + \sigma^2 I)^{-1} K(X, X_\star) \end{aligned}$$

Choice of Kernel Function

The expression for the posterior above shows that when the observed data has low covariance with the data we wish to predict (i.e. when $K(X_\star, X)$ is filled with small values compared to $K(X_\star, X_\star)$), the posterior will revert to the prior distribution specified by the GP. It is therefore important to correctly formulate the kernel function and mean depending on the data. For the mean, we assumed a constant mean of 0 after normalising the data. Choosing a kernel function must be done with some more care. The tide data is periodic, however there is also some variation in the height of the peaks and troughs. Therefore in this report the kernel function used is a combination of the periodic and RBF kernels, as below.

$$\begin{aligned} K(x, x') &= K_{\text{periodic}}(x, x') + \alpha K_{\text{rbf}}(x, x') \\ K_{\text{periodic}}(x, x') &= \sigma_p^2 \exp \left(-\frac{2(\sin(\pi d(x, x')/p))}{\omega_p} \right) \\ K_{\text{rbf}}(x, x') &= \sigma_r^2 \exp \left(-\frac{\|x - x'\|^2}{2\omega_r} \right) \end{aligned}$$

where $d(x, x')$ is the euclidean distance between two points x and x' . I found this combination made it easier to fit both the variation in height of the peaks and the periods of the function because it allows periodic and adjacent

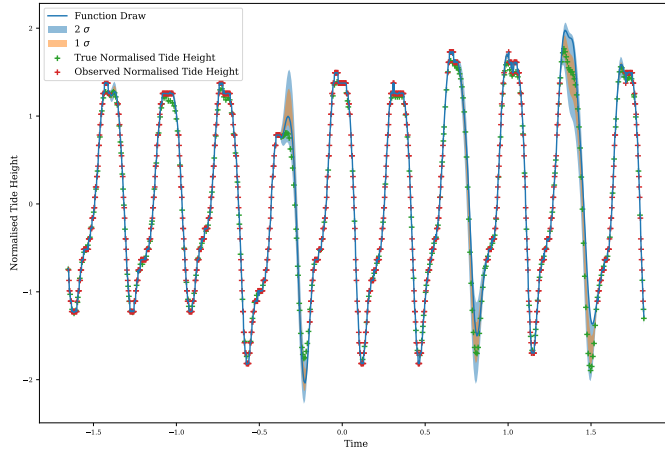


Figure 1: Plot showing the fit of the Gaussian Process to the normalised tide height

Figure 2: Animation showing randomly initialised models fitted to progressively larger portions of the dataset. This will animate when opened with Adobe Acrobat Reader.

similarity to be balanced. This kernel function has a number of hyperparameters, which we must optimise over. We do not optimise over the noise variance σ , but it was estimated to 1 significant figure from the data as 0.03.

Optimising Hyperparameters

To choose the hyperparameters, we do some optimisation. The quantity that we optimise is the marginal likelihood:

$$\begin{aligned} p(\mathbf{y}|X) &= \int p(\mathbf{y}|\mathbf{f}, X)p(\mathbf{f}|X)d\mathbf{f} \\ &= \mathcal{N}(\mathbf{0}, K(X, X) + \sigma^2 I) \end{aligned}$$

This quantity measures, for a particular set of hyperparameters, the probability density of the observations that we recorded across all possible function draws. This is different to the likelihood $p(\mathbf{y}|\mathbf{f}, X)$, which only considers one particular function drawn from the distribution. I optimised the log-marginal likelihood with the standard SciPy implementation of LBFGS-B.

Experiments

I first fit a Gaussian Process with the kernel specified above to the whole dataset. This is shown in Figure 1. This was obtained by specifying an initial guess for the hyperparameters chosen by hand, and the optimising them as above. The prediction of the mean on this data is clearly very well fit, but the uncertainty prediction could be better. Its form is good because the uncertainty is greater when there is a poverty of data, and low when nearby sensor data is available. Moreover, the uncertainty is higher when predicting the peaks and troughs than the data in between, which appears more regular. However, the overall uncertainty is unreasonably low, particularly in larger regions where there is little or no data.

Sequential Prediction

To explore the uncertainty estimates further, I looked at sequential predictions on the sensor data. The results are shown in the animation in Figure 2. This can be viewed in Adobe Acrobat. In this task I split the data into 10 chunks, and then fit a GP starting from random initial hyperparameters using the method above on progressively larger numbers of chunks. With the bounds I was using for the whole data set, random initialisations had very variable performance. I therefore used much tighter bounds with the random initial seeds to allow the GP to fit the data properly, while avoiding narrowing the search to a known set of good parameters.

For the first 10% of the data, the GP has not seen enough data to correctly estimate the period. It therefore guesses slightly incorrectly and is far too certain about the data in some sections. The model fit on 20% of the data is similar, although its uncertainty is much higher. In both cases there is likely not enough data for the GP to understand how uncertain it should be. If the model has not seen at least 1 period, it cannot possibly estimate the period, or the variation in the data across periods. The third model has correctly fit the period, but not correct uncertainty estimates. I think this is caused in part by there being little variation in the heights of the first three periods compared with the fourth, which is much lower. The fourth model does better, and the fifth is the first that truly fits the rest of the data reasonably, with none of the periods representing significant deviations from its expectations. After this, the differences between the models are largely due to random variation. The eighth model is particularly poor, which I believe is due to poor optimisation performance. The model before it had a maximum marginal likelihood of 1154, compared to its value of 1054. I would expect this to usually increase with more data.

A Supplementary Figures




Figure 3: Diagram illustrating a Gaussian distribution over function values at \mathbf{x}_* . Some simple functions drawn from the distribution are shown in dark green.

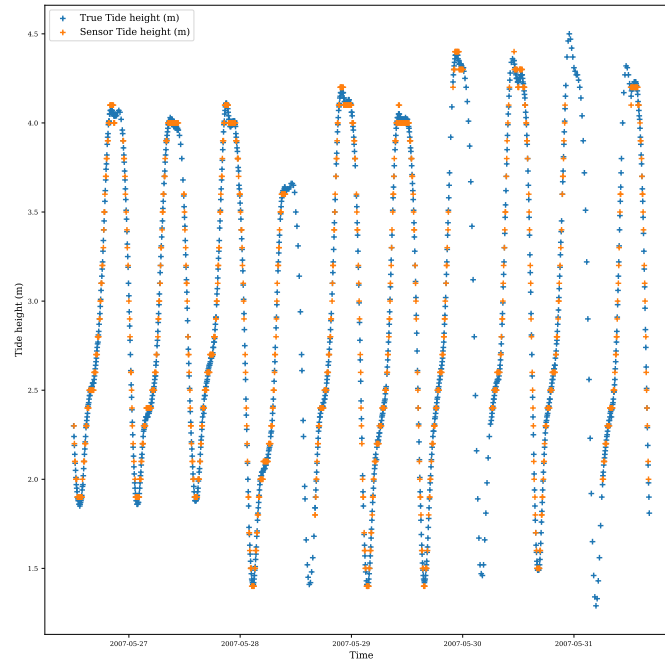


Figure 4: Scatter plot of the data obtained by the sensor and the true tide height